




Convolutional Neural Networks


Outline of Presentation

- ❖ Image and Computers
- ❖ Features Extraction
- ❖ Manual Feature Extraction
- ❖ Manual feature Representation
- ❖ Feature Extraction Issues
- ❖ Fully Connected Neural Network
- ❖ Feature Extraction and Convolution
- ❖ Feature Matching
- ❖ Convolution Operation
- ❖ Convolution Layer and Feature Maps
- ❖ Convolutional Neural Network: Spatial View
- ❖ Pooling and Activation Function (ReLU)
- ❖ CNNs Applications

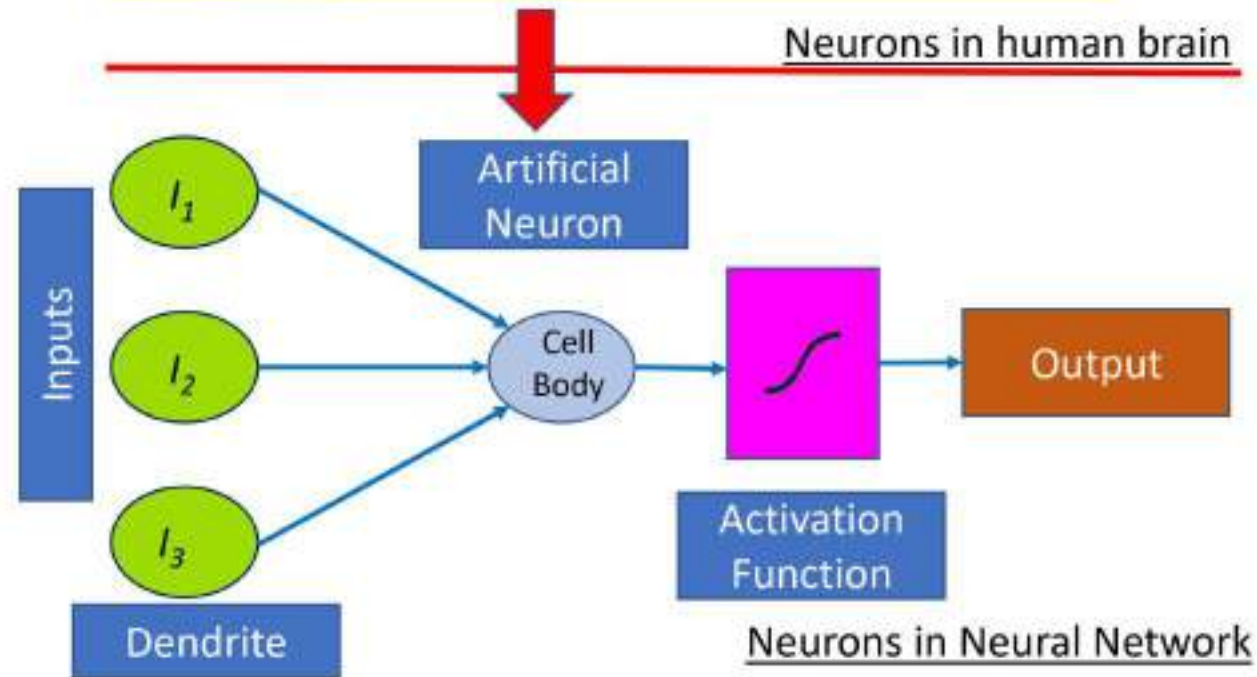
Visual Cortex and Deep Neural Network (DNN)



Visual Cortex and Deep Neural Network (DNN)



Neurons in human brain



Comparison of Biological neural network and Artificial Neural Network

Biological Neural Network	Artificial Neural Network
Stimulus	Input
Receptor	Input Layer
Neural Network	Processing Layer
Neuron	Processing Element
Dendrite	Addition Function
Synapse	Weight
Cell Body	Transfer/ Activation Function
Axon	Artificial Neural Output

Images are Just a Matrix of Numbers



255	200	211	235	1
200	161	217	233	0
218	65	214	237	0
232	29	217	236	1
234	23	216	240	0
102	31	217	234	0


255	200	211	235	1
200	161	217	233	0
218	65	214	237	0
232	29	217	236	0
234	23	216	240	0
102	31	217	234	0

Computer Interpretation

- For a grayscale images, the pixel value is a single number that represents the brightness of the pixel. The most common pixel format is the byte image, where this number is stored as an 8-bit integer giving a range of possible values from 0 to 255.
- Similarly for color images, each level is represented by the range of decimal numbers from 0 to 255 (256 levels for each color), equivalent to the range of binary numbers from 00000000 to 11111111, or hexadecimal 00 to FF. The total number of available colors is 256 x 256 x 256, or 16,777,216 possible color.

Computer Vision Task

- **Regression** - is about predicting a quantity. It is the process of finding a model or function for distinguishing the data into continuous real values. Example: Price estimation of houses.
- **Classification**- is about predicting a label. It is the process of finding or discovering a model or function which helps in separating the data into multiple categorical classes i.e. discrete values.



High Level Feature Detection

Low-level features are minor details of the image, like lines or dots, that can be picked up by convolutional filter (for really low-level things) or SIFT or HOG (for more abstract things like edges). High-level features are built on top of low-level features to detect objects and larger shapes in the image.



Manual Feature Extraction

1. *Domain Knowledge*



2. Specify Features



3. Detect Features for Classification



Manual Feature Extraction: Issues



Occlusion



Cluttered Background



Intra-class Diversity



Object Deformation



Scale Variation



Illumination Variation

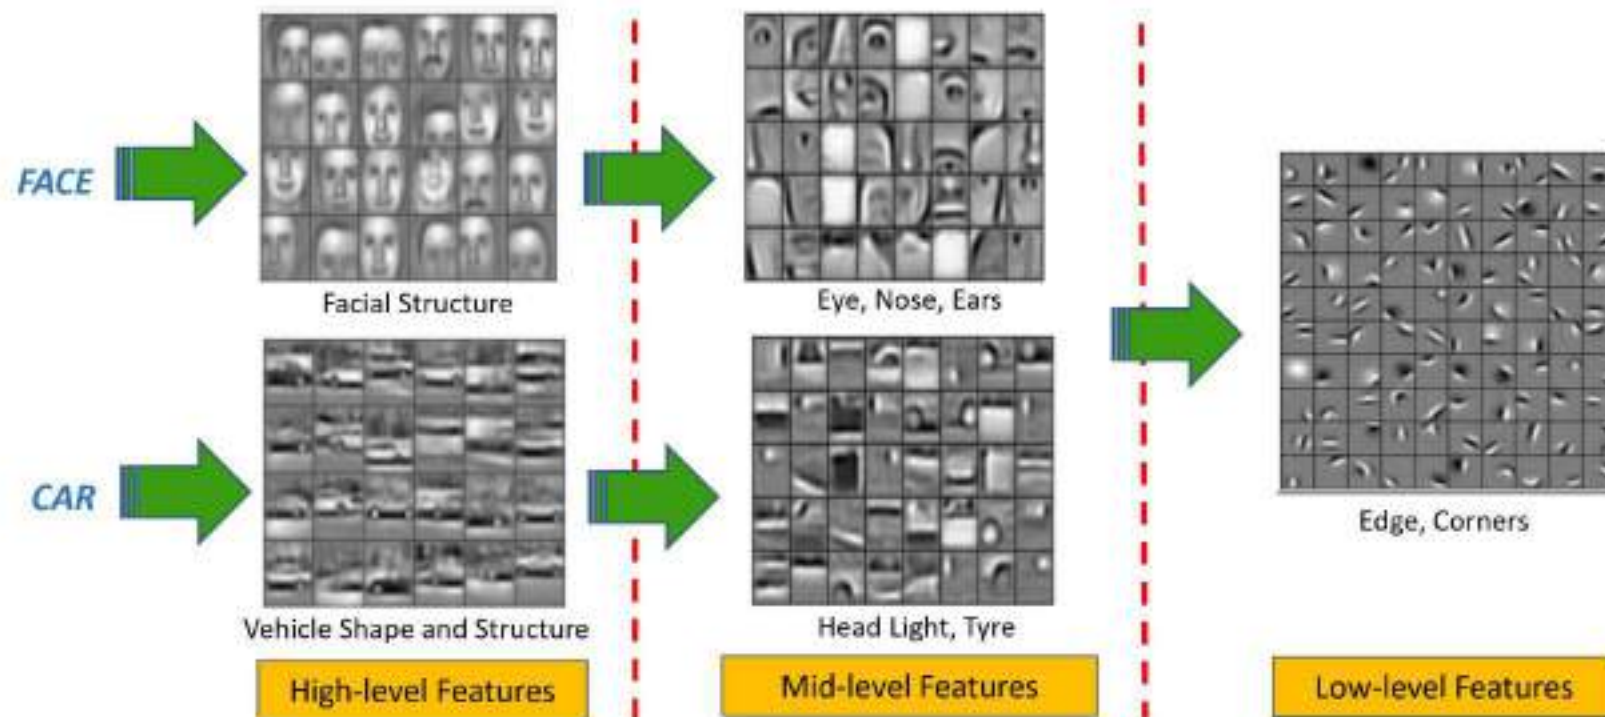


Change in Viewing Angles




Image Deformity

Feature Representation

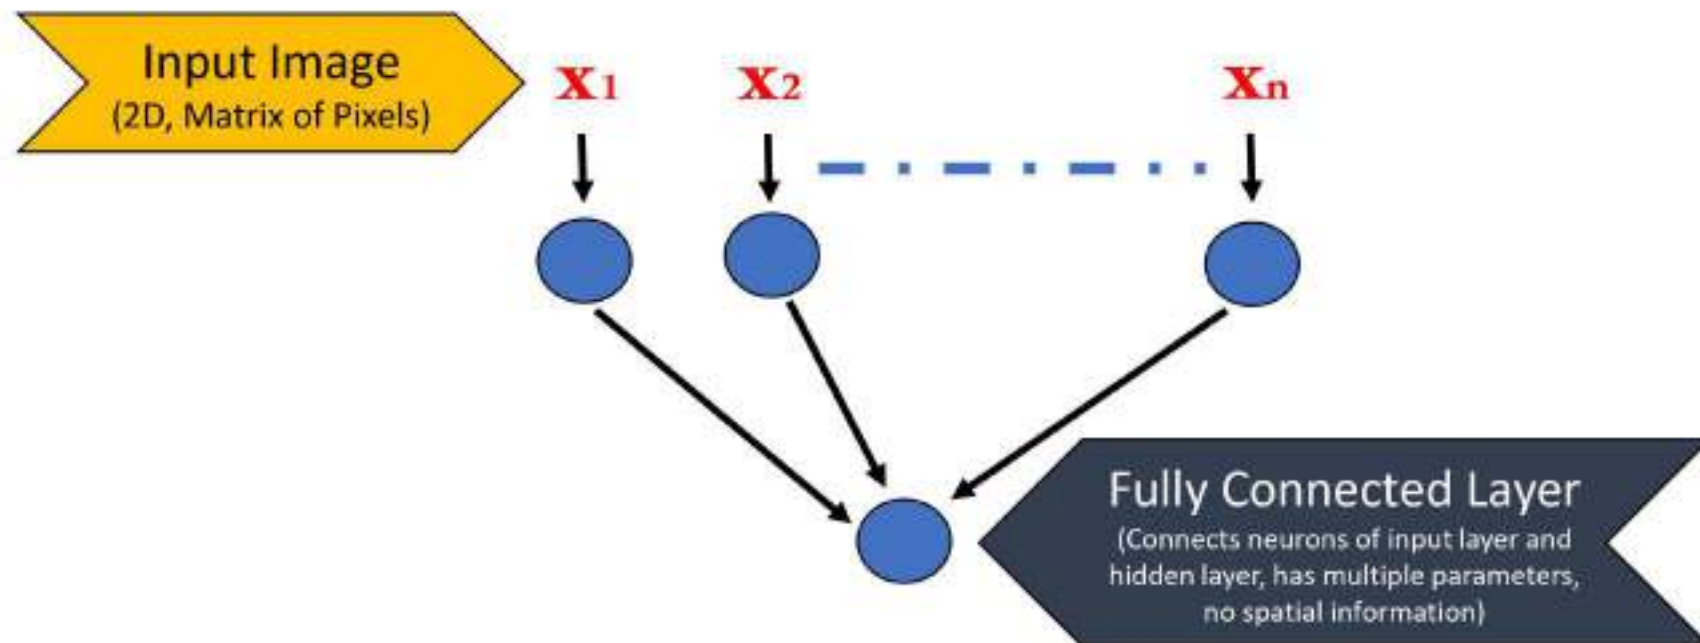


Fully Connected Neural Network




Fully connected neural networks (FCNNs) are a type of artificial neural network where the architecture is such that all the nodes, or neurons, in one layer are connected to the neurons in the next layer.

Fully Connected Neural Network



Feature Extraction using Convolution



Connect Input Layer patches to neurons of hidden layer/subsequent layer with sliding window approach.

Step 1:

Extract Set of Local Features by applying filters (set of weights)

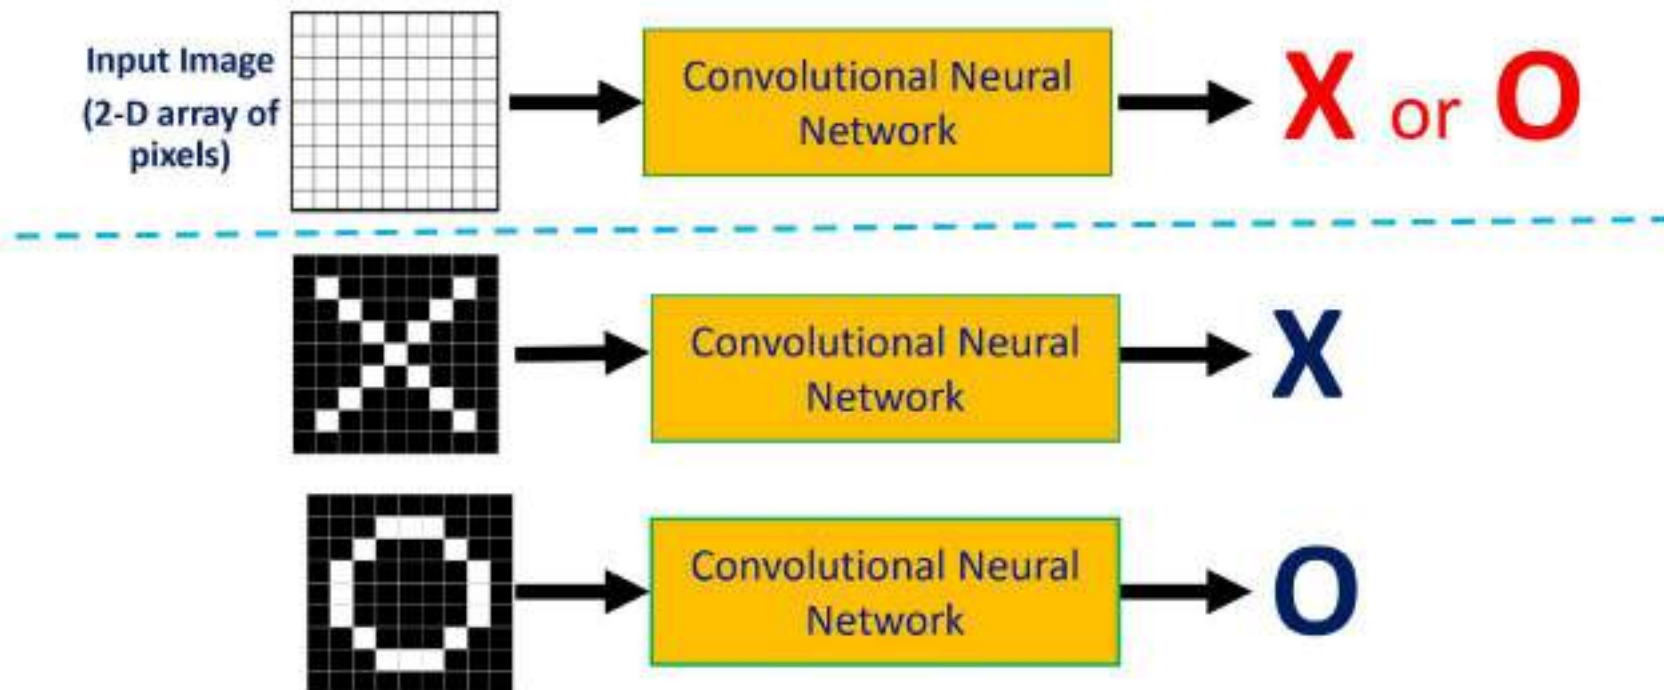
Step 2:

Apply Multiple Filters for extraction of different features

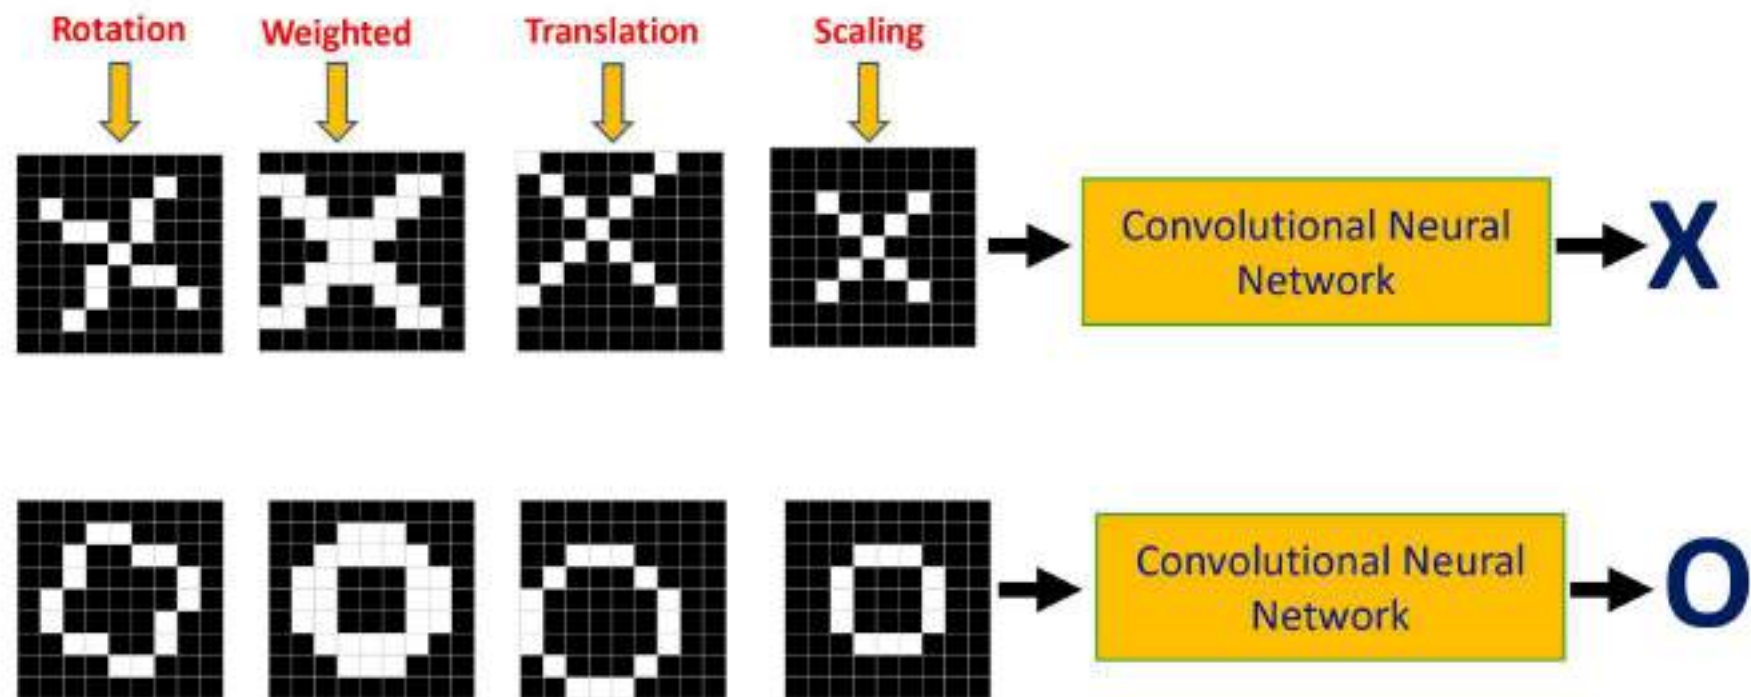
Step 3:

Spatial Sharing of parameters for each filter


Convolutional Neural Network (Feature Extraction and Convolution)



Challenging Cases



Computer and Human Interpretation



Computer Interpretation



-1	-1	-1	-1	-1	-1	-1	-1
-1	1	-1	-1	-1	-1	-1	-1
-1	-1	1	-1	-1	-1	1	-1
-1	-1	-1	1	-1	1	-1	-1
-1	-1	-1	-1	1	-1	-1	-1
-1	-1	-1	-1	-1	1	-1	-1
-1	-1	1	-1	-1	-1	1	-1
-1	1	-1	-1	-1	-1	-1	1
-1	-1	-1	-1	-1	-1	-1	-1

Computer Interpretation

Pixel wise
Matching



-1	-1	-1	-1	-1	-1	-1	-1	-1
-1	X	-1	-1	-1	-1	X	X	-1
-1	X	X	-1	-1	X	X	-1	-1
-1	-1	X	1	-1	1	-1	-1	-1
-1	-1	-1	-1	1	-1	-1	-1	-1
-1	-1	-1	1	-1	1	X	-1	-1
-1	-1	X	X	-1	-1	X	X	-1
-1	X	X	-1	-1	-1	-1	X	-1
-1	-1	-1	-1	-1	-1	-1	-1	-1

Decision




-1	-1	-1	-1	-1	-1	-1	-1	-1
-1	1	-1	-1	-1	-1	-1	1	-1
-1	-1	1	-1	-1	-1	1	-1	-1
-1	-1	-1	1	-1	1	-1	-1	-1
-1	-1	-1	-1	1	-1	-1	-1	-1
-1	-1	-1	1	-1	1	-1	-1	-1
-1	-1	1	-1	-1	-1	1	-1	-1
-1	1	-1	-1	-1	-1	-1	1	-1
-1	-1	-1	-1	-1	-1	-1	-1	-1





-1	-1	-1	-1	-1	-1	-1	-1	-1
-1	-1	-1	-1	-1	-1	1	-1	-1
-1	1	-1	-1	-1	1	-1	-1	-1
-1	-1	1	1	-1	1	-1	-1	-1
-1	-1	-1	-1	1	-1	-1	-1	-1
-1	-1	-1	1	-1	1	1	-1	-1
-1	-1	-1	1	-1	-1	-1	1	-1
-1	-1	1	-1	-1	-1	-1	-1	-1
-1	-1	-1	-1	-1	-1	-1	-1	-1

Computers are Literal


Feature matching for symbol 'X'



Piece Matching of Features



Piece Matching of Features



Convolution Operation

1 <small>x1</small>	1 <small>x0</small>	1 <small>x1</small>	0	0
0 <small>x0</small>	1 <small>x1</small>	1 <small>x0</small>	1	0
0 <small>x1</small>	0 <small>x0</small>	1 <small>x1</small>	1	1
0	0	1	1	0
0	1	1	0	0

Image

4		

Convolved
Feature

Convolution Operation

1	-1	-1
-1	1	-1
-1	-1	1

$$\boxed{1} \times \boxed{1} = 1$$

[illegible]

Convolution Operation

1	-1	-1
-1	1	-1
-1	-1	1

$$\boxed{1} \times \boxed{1} = 1$$

-1	-1	-1	-1	-1	-1	-1	-1	-1
-1	1	-1	-1	-1	-1	-1	1	-1
-1	-1	1	-1	-1	-1	1	-1	-1
-1	-1	-1	1	-1	1	-1	-1	-1
-1	-1	-1	-1	1	-1	-1	-1	-1
-1	-1	-1	1	-1	1	-1	-1	-1
-1	-1	1	-1	-1	-1	1	-1	-1
-1	1	-1	-1	-1	-1	-1	1	-1
-1	-1	-1	-1	-1	-1	-1	-1	-1

1		

Convolution Operation

1	-1	-1
-1	1	-1
-1	-1	1

$$\begin{array}{|c|} \hline -1 \\ \hline \end{array} \times \begin{array}{|c|} \hline -1 \\ \hline \end{array} = 1$$

-1	-1	-1	-1	-1	-1	-1	-1	-1
-1	1	-1	-1	-1	-1	-1	1	-1
-1	-1	1	-1	-1	-1	1	-1	-1
-1	-1	-1	1	-1	1	-1	-1	-1
-1	-1	-1	-1	1	-1	-1	-1	-1
-1	-1	-1	1	-1	1	-1	-1	-1
-1	-1	1	-1	-1	-1	1	-1	-1
-1	1	-1	-1	-1	-1	-1	1	-1
-1	-1	-1	-1	-1	-1	-1	-1	-1

1	1	

Convolution Operation

1	-1	-1
-1	1	-1
-1	-1	1

$$\begin{array}{|c|} \hline -1 \\ \hline \end{array} \times \begin{array}{|c|} \hline -1 \\ \hline \end{array} = 1$$

-1	-1	-1	-1	-1	-1	-1	-1	-1
-1	1	-1	-1	-1	-1	-1	1	-1
-1	-1	1	-1	-1	-1	1	-1	-1
-1	-1	-1	1	-1	1	-1	-1	-1
-1	-1	-1	-1	1	-1	-1	-1	-1
-1	-1	-1	1	-1	1	-1	-1	-1
-1	-1	1	-1	-1	-1	1	-1	-1
-1	1	-1	-1	-1	-1	-1	1	-1
-1	-1	-1	-1	-1	-1	-1	-1	-1

1	1	1

Convolution Operation

1	-1	-1
-1	1	-1
-1	-1	1

$$\boxed{-1} \times \boxed{-1} = 1$$

-1	-1	-1	-1	-1	-1	-1	-1	-1
-1	1	-1	-1	-1	-1	-1	1	-1
-1	-1	1	-1	-1	-1	1	-1	-1
-1	-1	-1	1	-1	1	-1	-1	-1
-1	-1	-1	-1	1	-1	-1	-1	-1
-1	-1	-1	1	-1	1	-1	-1	-1
-1	-1	1	-1	-1	-1	1	-1	-1
-1	1	-1	-1	-1	-1	-1	1	-1
-1	-1	-1	-1	-1	-1	-1	-1	-1

1	1	1
1		

Convolution Operation

1	-1	-1
-1	1	-1
-1	-1	1

$$\boxed{1} \times \boxed{1} = 1$$

-1	-1	-1	-1	-1	-1	-1	-1	-1
-1	1	-1	-1	-1	-1	-1	1	-1
-1	-1	1	-1	-1	-1	1	-1	-1
-1	-1	-1	1	-1	1	-1	-1	-1
-1	-1	-1	-1	1	-1	-1	-1	-1
-1	-1	-1	1	-1	1	-1	-1	-1
-1	-1	1	-1	-1	-1	1	-1	-1
-1	1	-1	-1	-1	-1	-1	1	-1
-1	-1	-1	-1	-1	-1	-1	-1	-1

1	1	1
1	1	

Convolution Operation

1	-1	-1
-1	1	-1
-1	-1	1

$$\boxed{-1} \times \boxed{-1} = 1$$

-1	-1	-1	-1	-1	-1	-1	-1	-1
-1	1	-1	-1	-1	-1	-1	1	-1
-1	-1	1	-1	-1	-1	1	-1	-1
-1	-1	-1	1	-1	1	-1	-1	-1
-1	-1	-1	-1	1	-1	-1	-1	-1
-1	-1	-1	1	-1	1	-1	-1	-1
-1	-1	1	-1	-1	-1	1	-1	-1
-1	1	-1	-1	-1	-1	-1	1	-1
-1	-1	-1	-1	-1	-1	-1	-1	-1

1	1	1
1	1	1

Convolution Operation

1	-1	-1
-1	1	-1
-1	-1	1

$$\boxed{-1} \times \boxed{-1} = 1$$

-1	-1	-1	-1	-1	-1	-1	-1	-1
-1	1	-1	-1	-1	-1	-1	1	-1
-1	-1	1	-1	-1	-1	1	-1	-1
-1	-1	-1	1	-1	1	-1	-1	-1
-1	-1	-1	-1	1	-1	-1	-1	-1
-1	-1	-1	1	-1	1	-1	-1	-1
-1	-1	1	-1	-1	-1	1	-1	-1
-1	1	-1	-1	-1	-1	-1	1	-1
-1	-1	-1	-1	-1	-1	-1	-1	-1

1	1	1
1	1	1
1		

Convolution Operation

1	-1	-1
-1	1	-1
-1	-1	1

$$\boxed{-1} \times \boxed{-1} = 1$$

-1	-1	-1	-1	-1	-1	-1	-1	-1
-1	1	-1	-1	-1	-1	-1	1	-1
-1	-1	1	-1	-1	-1	1	-1	-1
-1	-1	-1	1	-1	1	-1	-1	-1
-1	-1	-1	-1	1	-1	-1	-1	-1
-1	-1	-1	1	-1	1	-1	-1	-1
-1	-1	1	-1	-1	-1	1	-1	-1
-1	1	-1	-1	-1	-1	-1	1	-1
-1	-1	-1	-1	-1	-1	-1	-1	-1

1	1	1
1	1	1
1	1	

Convolution Operation


1	-1	-1
-1	1	-1
-1	-1	1

$$\boxed{1} \times \boxed{1} = 1$$


-1	-1	-1	-1	-1	-1	-1	-1	-1
-1	1	-1	-1	-1	-1	-1	1	-1
-1	-1	1	-1	-1	-1	1	-1	-1
-1	-1	-1	1	-1	1	-1	-1	-1
-1	-1	-1	-1	1	-1	-1	-1	-1
-1	-1	-1	1	-1	1	-1	-1	-1
-1	-1	1	-1	-1	-1	1	-1	-1
-1	1	-1	-1	-1	-1	-1	1	-1
-1	-1	-1	-1	-1	-1	-1	-1	-1

1	1	1
1	1	1
1	1	1

Convolution Operation




Convolution Operation



Convolution Operation

-1	-1	-1	-1	-1	-1	-1	-1	-1
-1	1	-1	-1	-1	-1	-1	1	-1
-1	-1	1	-1	-1	-1	1	-1	-1
-1	-1	-1	1	-1	1	-1	-1	-1
-1	-1	-1	-1	1	-1	-1	-1	-1
-1	-1	-1	1	-1	1	-1	-1	-1
-1	-1	1	-1	-1	-1	1	-1	-1
-1	1	-1	-1	-1	-1	-1	1	-1
-1	-1	-1	-1	-1	-1	-1	-1	-1




1	-1	-1
-1	1	-1
-1	-1	1


 $=$

0.77	-0.11	0.11	0.33	0.55	-0.11	0.33
-0.11	1.00	-0.11	0.33	-0.11	0.11	-0.11
0.11	-0.11	1.00	-0.33	0.11	-0.11	0.55
0.33	0.33	-0.33	0.55	-0.33	0.33	0.33
0.55	-0.11	0.11	-0.33	1.00	-0.11	0.11
-0.11	0.11	-0.11	0.33	-0.11	1.00	-0.11
0.33	-0.11	0.55	0.33	0.11	-0.11	0.77


Convolution Operation



Convolutional layer



Convolutional layer



Feature Maps on an Image



Original Image



Sharpen



Edge Detected

Kernels and their effects on an Image



Original Image



0	0	0
0	1	0
0	0	0

Identity Kernel



Output Image – Same as Original

Kernels and their effects on an Image



Original Image



0.0625	0.125	0.0625
0.125	0.25	0.125
0.0625	0.125	0.0625

Blur



Output Image

Kernels and their effects on an Image



Original Image



1	0	-1
2	0	-2
1	0	-1

Left Sobel



Output Image

Sobel kernels are used to show *only* the differences in adjacent pixel values in a particular direction

Kernels and their effects on an Image



Original Image



-1	0	1
-2	0	2
-1	0	1

Right Sobel



Output Image

Sobel kernels are used to show *only* the differences in adjacent pixel values in a particular direction

Kernels and their effects on an Image



Original Image



-1	-2	-1
0	0	0
1	2	1

Bottom Sobel



Output Image

Sobel kernels are used to show *only* the differences in adjacent pixel values in a particular direction

Kernels and their effects on an Image



Original Image



1	2	1
0	0	0
-1	-2	-1

Top Sobel



Output Image

Sobel kernels are used to show *only* the differences in adjacent pixel values in a particular direction

Kernels and their effects on an Image



Original Image



-2	-1	0
-1	1	1
0	1	2

Emboss



Output Image

The **emboss** kernel (similar to the **Sobel** kernel and sometimes referred to mean the same) gives the illusion of depth by emphasizing the differences of pixels in a given direction. In this case, in a direction along a line from the top left to the bottom right.

Kernels and their effects on an Image



Original Image



-1	-1	-1
-1	8	-1
-1	-1	-1

Outline



Output Image

An **outline** kernel (also called an "edge" kernel) is used to highlight large differences in pixel values. A pixel next to neighbor pixels with close to the same intensity will appear black in the new image while one next to neighbor pixels that differ strongly will appear white.

Kernels and their effects on an Image



Original Image



0	-1	0
-1	5	-1
0	-1	0


Sharpen




Output Image

The **sharpen** kernel emphasizes differences in adjacent pixel values. This makes the image look more vivid.

CONVOLUTIONAL NEURAL NETWORK--- CLASSIFICATION



Feature Extraction using Spatial Structure



Connect Input Layer patches to neurons of hidden layer/subsequent layer with sliding window approach.

Step 1:

Take inputs from the window (set of weights)


Step 2:

Calculation of weighted sum

Step 3:

Apply Bias

Feature Extraction using Spatial Structure



Step 1:

Take inputs from the window
(set of weights)

Step 2:

Calculation of weighted sum

Step 3:


Apply Bias

For (p, q) neuron in hidden layer

$$\sum_{i=1}^4 \sum_{j=1}^4 w_{ij} x_{i+p, j+q} + b$$


Where, w_{ij} = weights in matrix

Convolutional Neural Network--- Spatial View



Stride = Step size of filter, Receptive Field = Location of connected path in an input image

Non-Linearity in Convolutional Neural Network




Applied after every convolutional layer. The rectified linear activation (ReLU) function is a simple calculation that returns the value provided as input directly, or the value 0.0 if the input is 0.0 or less.

$$g(x) = \max(0, x)$$


Rectified Linear Units (ReLUs)




Rectified Linear Units (ReLUs)



Rectified Linear Units (ReLUs)



Rectified Linear Units (ReLU)



ReLU layer

A stack of images becomes a stack of images with no negative values.



POOLING

- *Dimensionality Reduction*
- *Preserve Spatial Invariance*

STEPS

1. Pick a window size (usually 2 or 3).
2. Pick a stride (usually 2).
3. Walk your window across your filtered images.
4. From each window, take the maximum value.


Pooling: Shrinking the Image stack

1. Pick a window size (usually 2 or 3).
2. Pick a stride (usually 2).
3. Walk your window across your filtered images.
4. From each window, take the maximum value.


The three types of pooling operations are:

- ❑ **Max pooling:** The maximum pixel value of the batch is selected.
- ❑ **Min pooling:** The minimum pixel value of the batch is selected.
- ❑ **Average pooling:** The average value of all the pixels in the batch is selected.


Max-Pooling




Max-Pooling




Max-Pooling




Max-Pooling



Max-Pooling



Max-Pooling



Max-Pooling

0.77	0.11	0.11	0.94	0.05	0.15	0.01
0.11	0.00	0.11	0.00	0.10	0.10	0.11
0.11	0.11	1.00	0.00	0.10	0.10	0.01
0.10	0.10	0.11	0.00	0.10	0.10	0.11
0.00	0.11	0.11	0.01	1.00	0.10	0.10
0.11	0.10	0.11	0.10	0.10	1.00	0.11
0.10	0.11	0.10	0.10	0.11	0.10	0.77



1.00	0.33	0.55	0.33
0.33	1.00	0.33	0.55
0.55	0.33	1.00	0.11
0.33	0.55	0.11	0.77

0.95	0.05	0.10	0.10	0.10	0.95	0.91
0.05	0.95	0.05	0.01	0.05	0.05	0.05
0.10	0.05	0.05	0.77	0.05	0.05	0.11
0.10	0.01	0.77	1.00	0.77	0.10	0.11
0.10	0.05	0.05	0.77	0.05	0.05	0.11
0.05	0.05	0.05	0.01	0.05	0.05	0.05
0.10	0.05	0.10	0.10	0.10	0.05	0.11



0.55	0.33	0.55	0.33
0.33	1.00	0.55	0.11
0.55	0.55	0.55	0.11
0.33	0.11	0.11	0.33


0.95	0.11	0.00	0.00	0.11	0.11	0.77
0.11	0.11	0.11	0.10	0.11	1.00	0.11
0.00	0.11	0.11	0.00	1.00	0.11	0.11
0.00	0.10	0.11	0.05	0.11	0.10	0.11
0.11	0.11	1.00	0.10	0.11	0.11	0.00
0.11	1.00	0.11	0.00	0.11	0.11	0.11
0.77	0.11	0.11	0.00	0.00	0.11	0.11



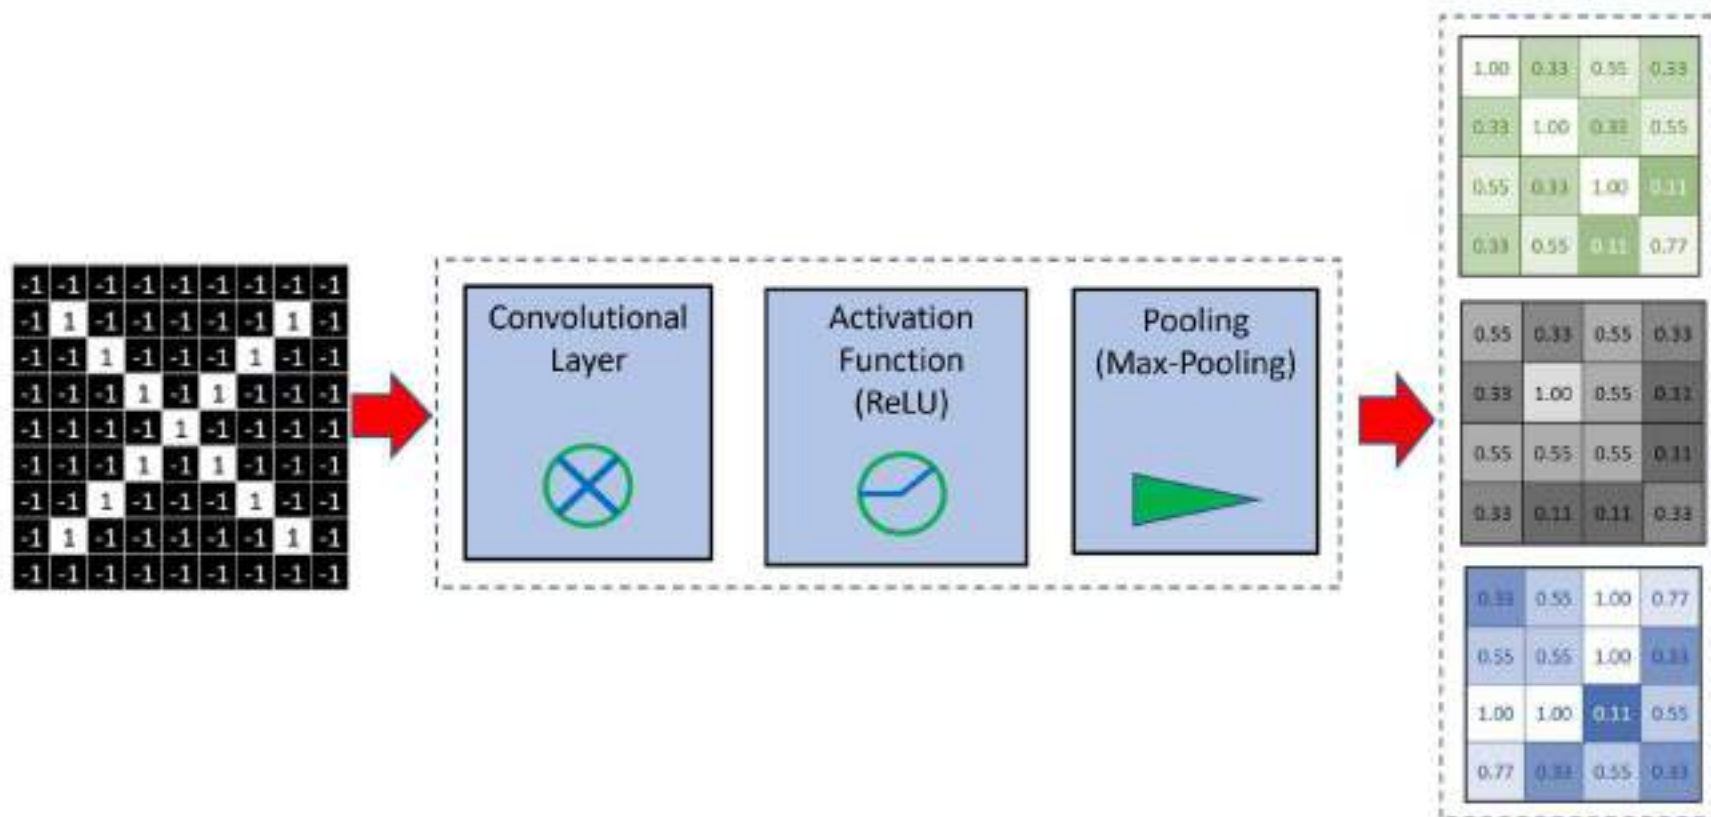
0.11	0.55	1.00	0.77
0.55	0.55	1.00	0.33
1.00	1.00	0.11	0.55
0.77	0.33	0.55	0.33

Pooling layer

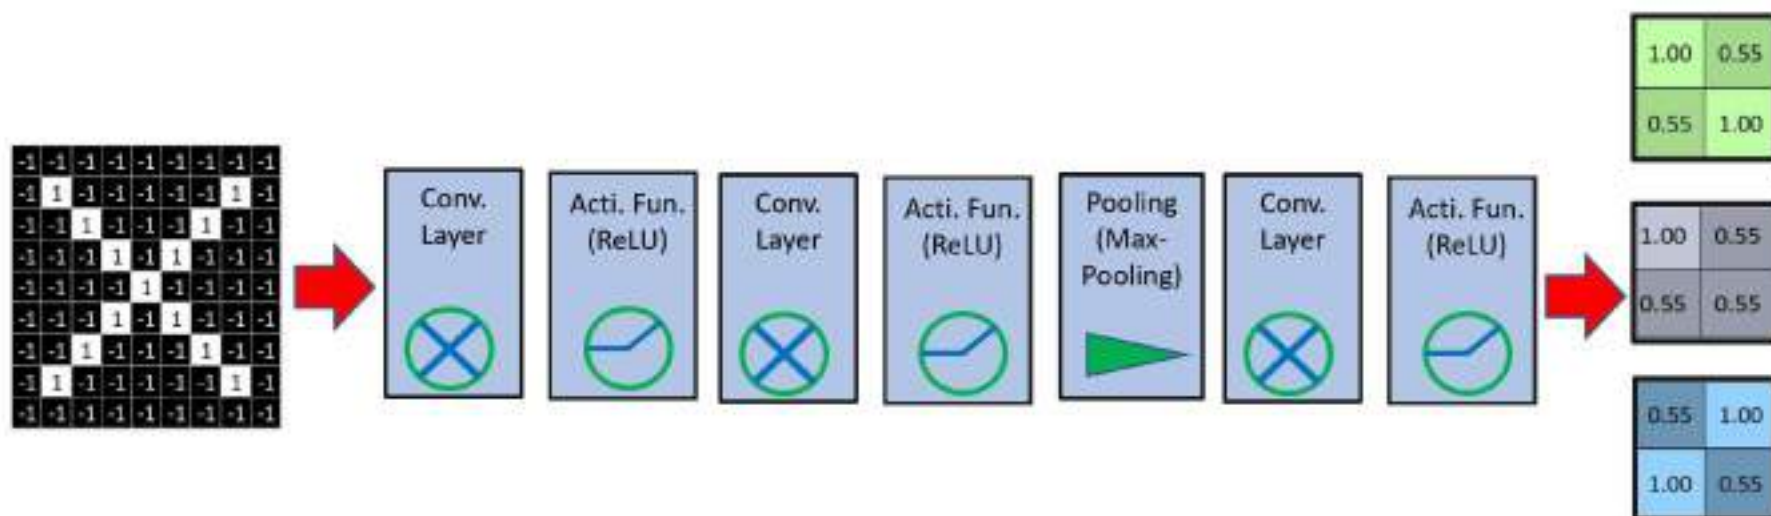
A stack of images becomes a stack of smaller images.




Stacking of Layers



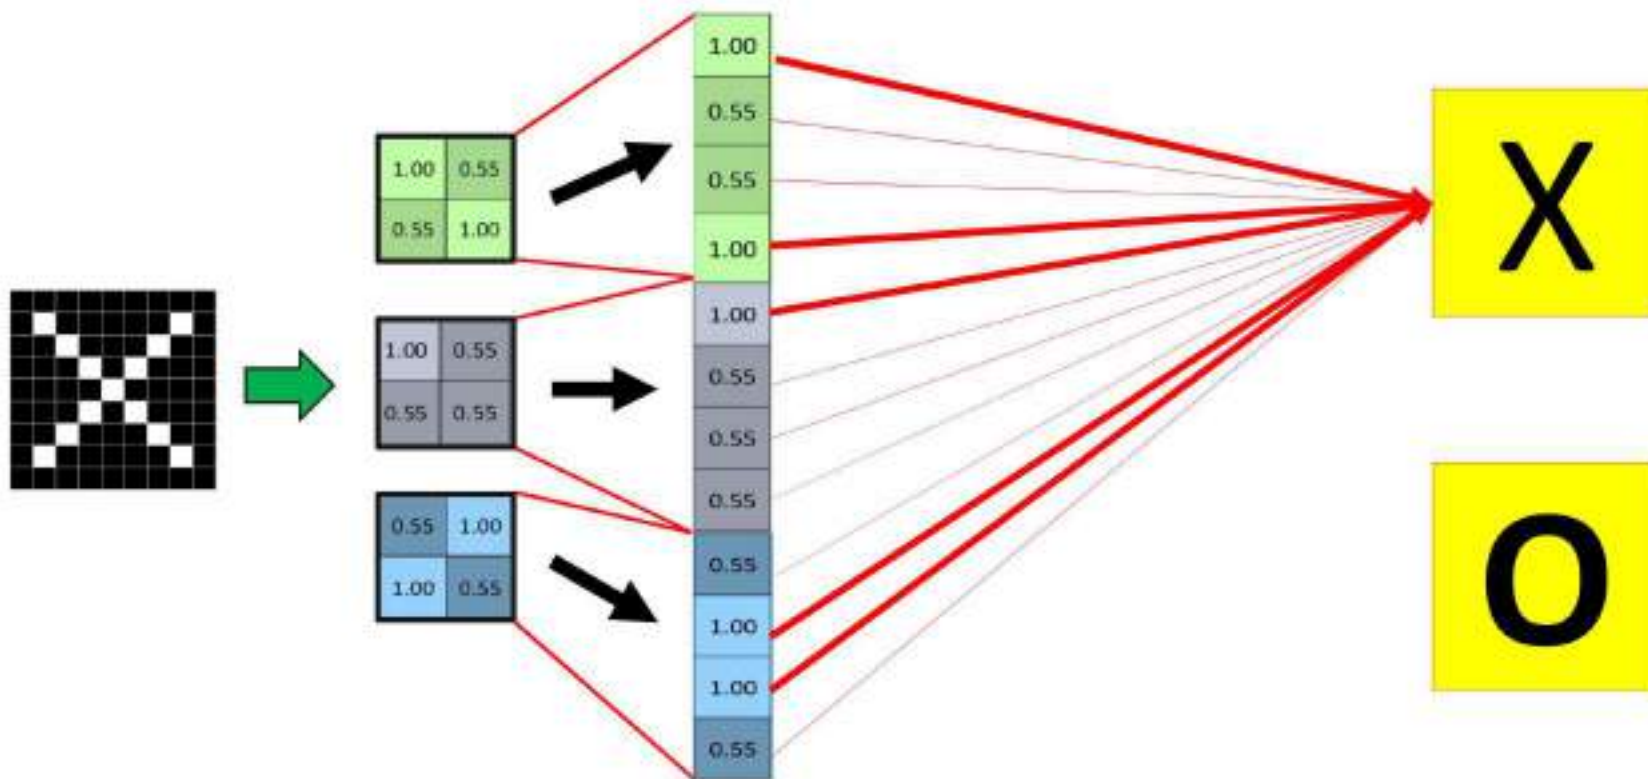
Multiple Stacking of Layers



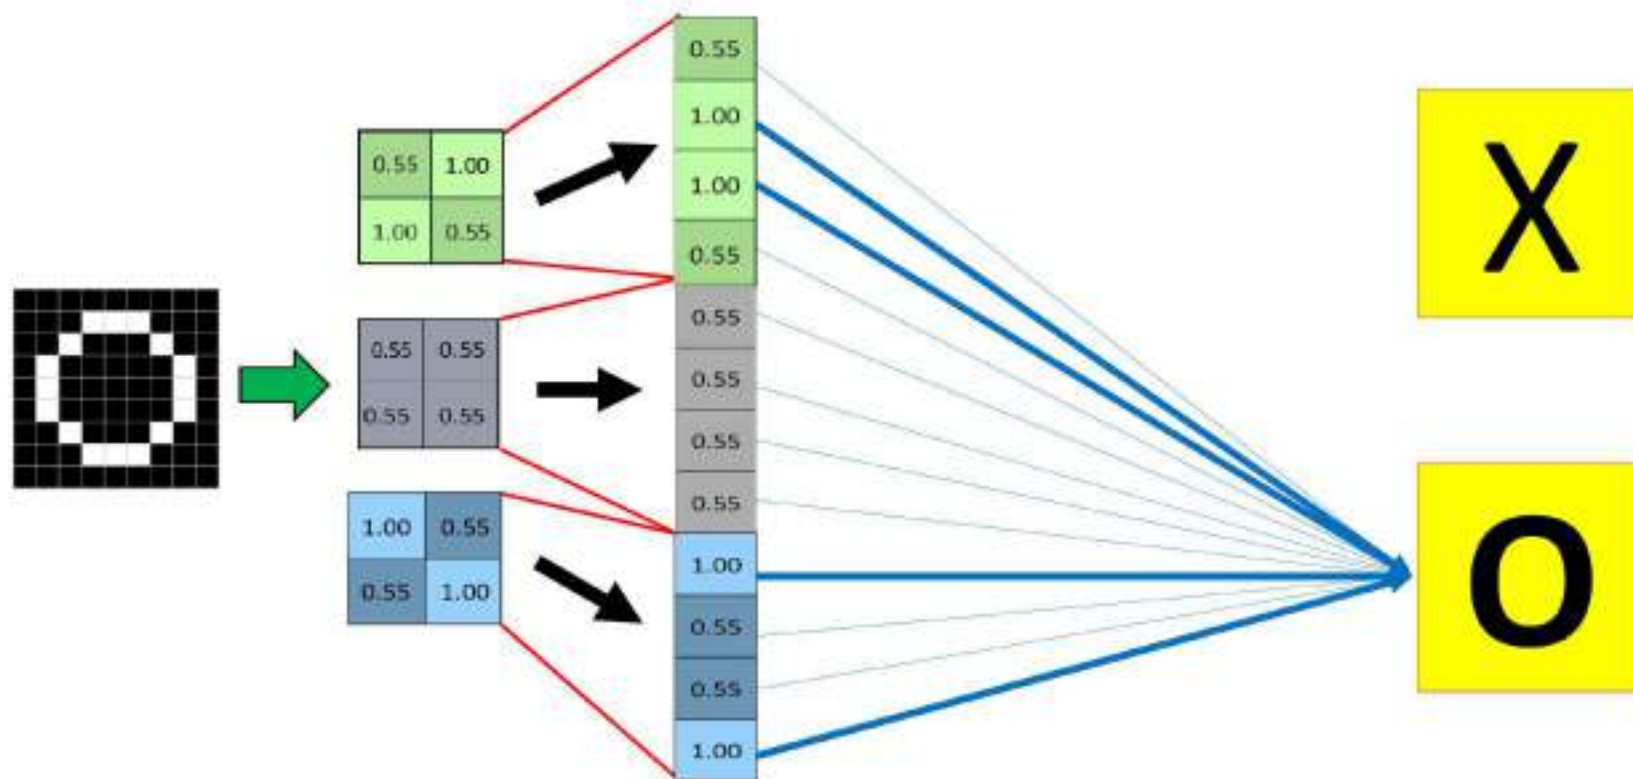
Fully Connected Layer (Training Phase)




Fully Connected Layer (Training Phase)



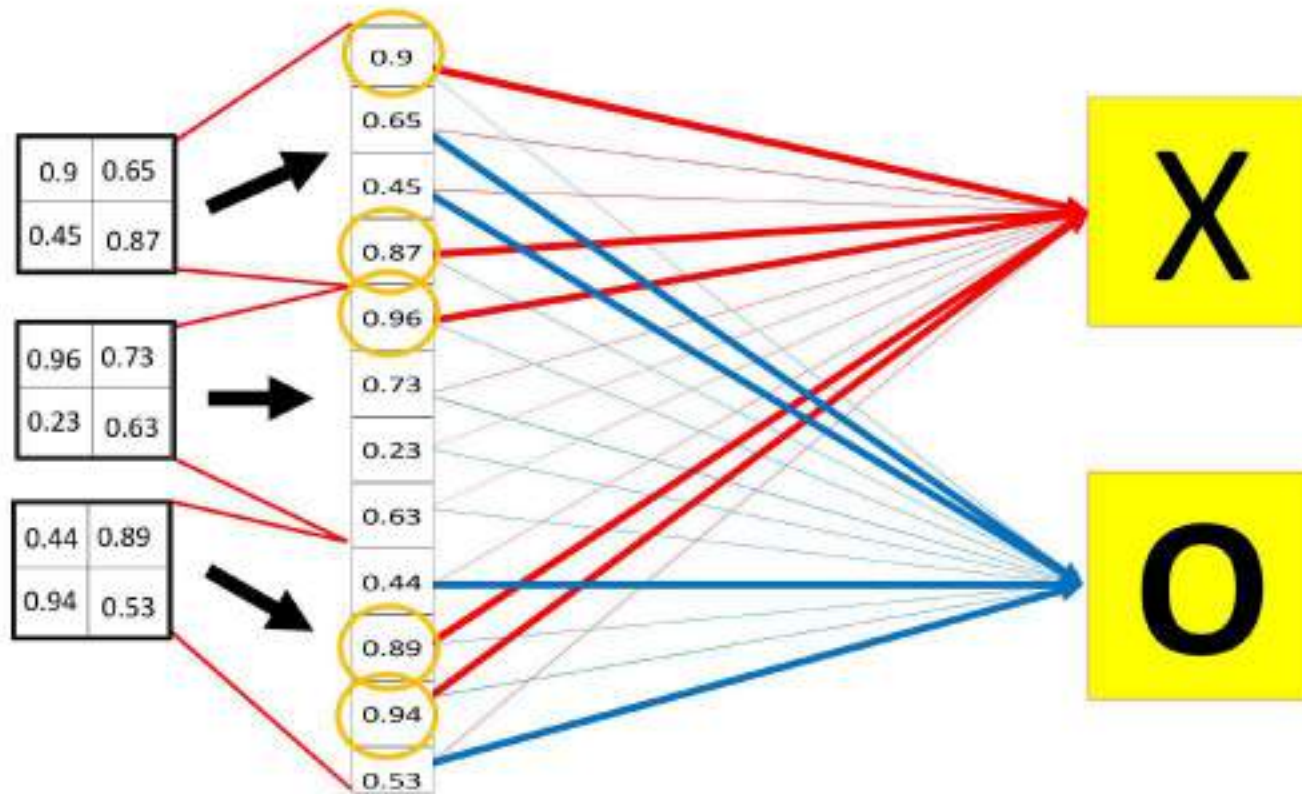
Fully Connected Layer (Training Phase)




Fully Connected Layer (Training Phase)



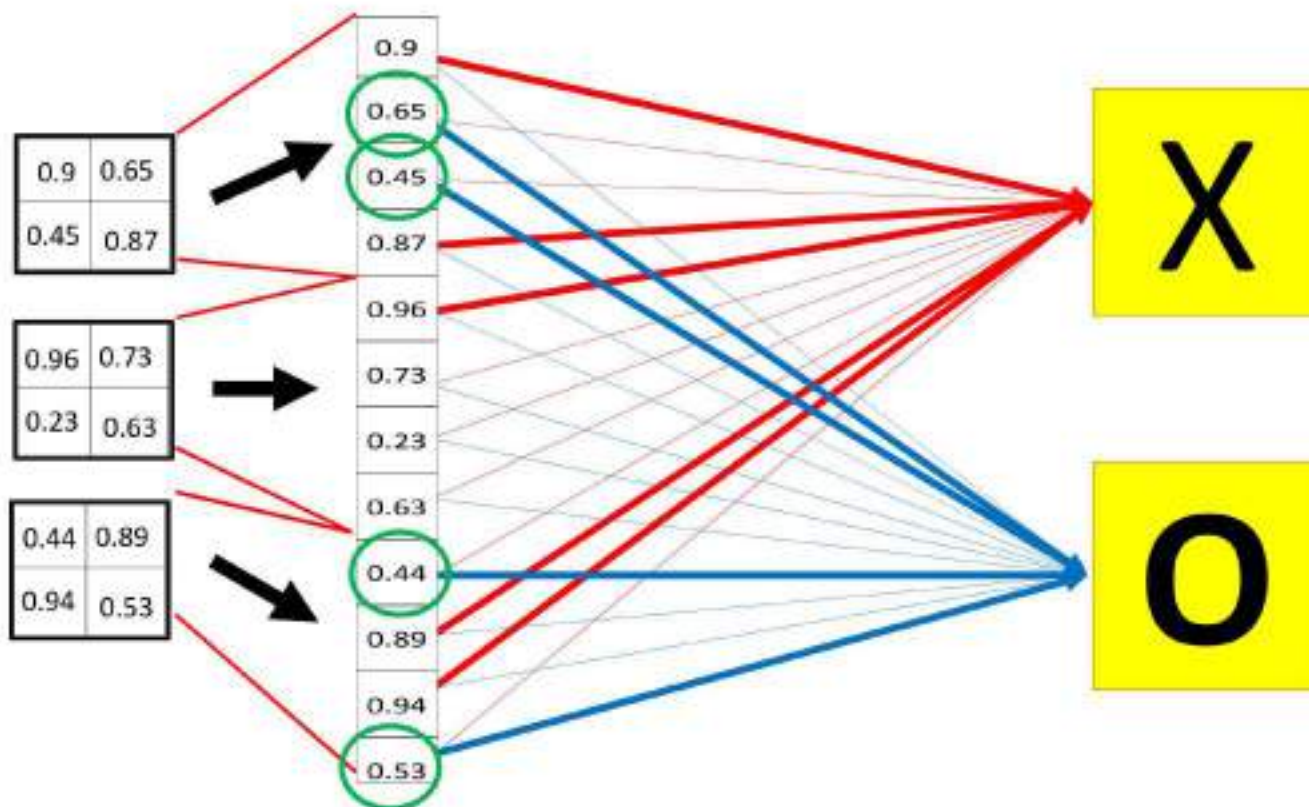
Fully Connected Layer (Testing Phase)



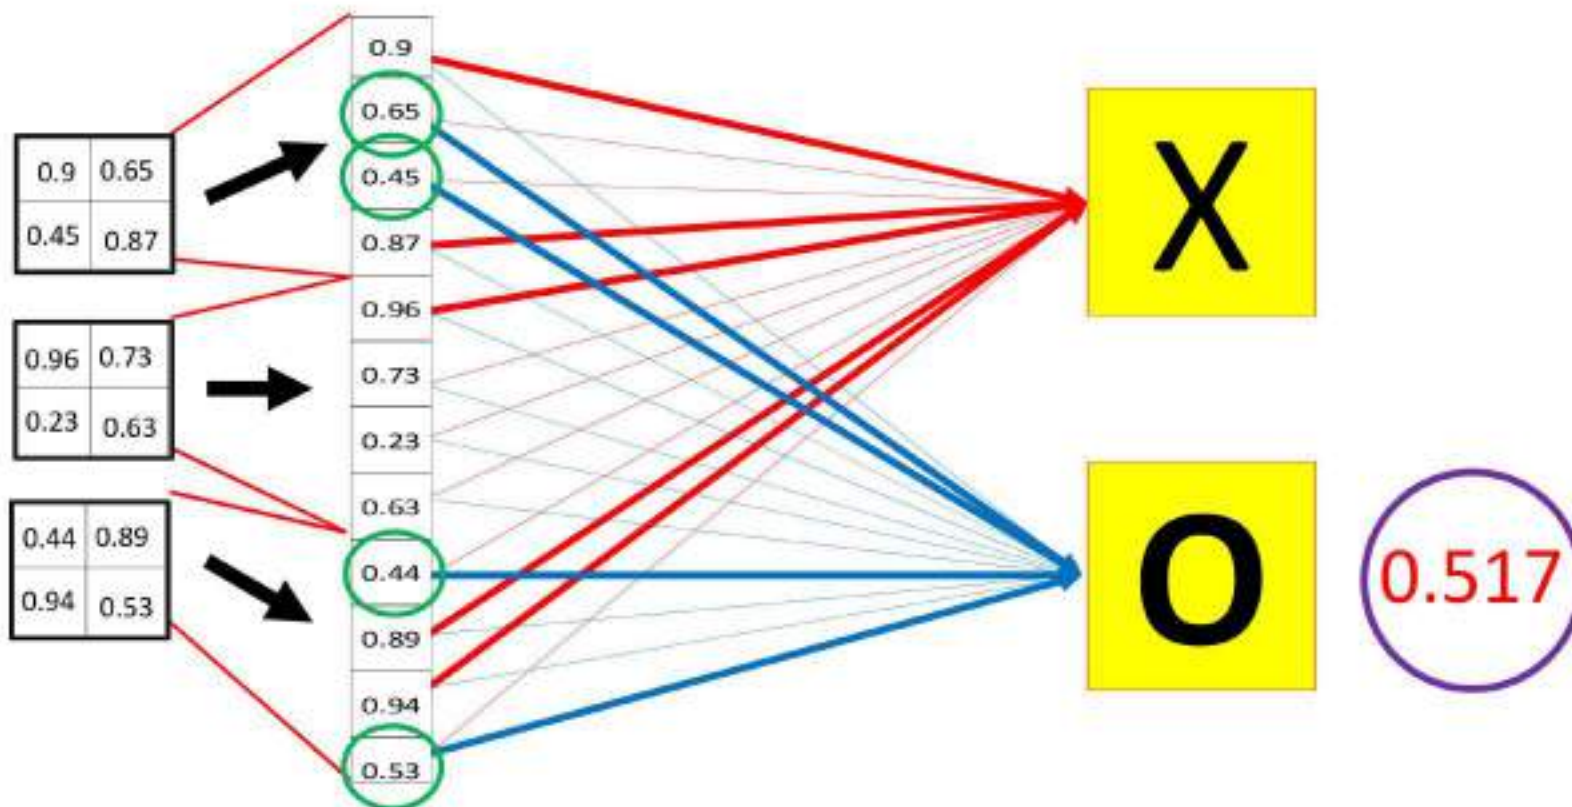
Fully Connected Layer (Testing Phase)



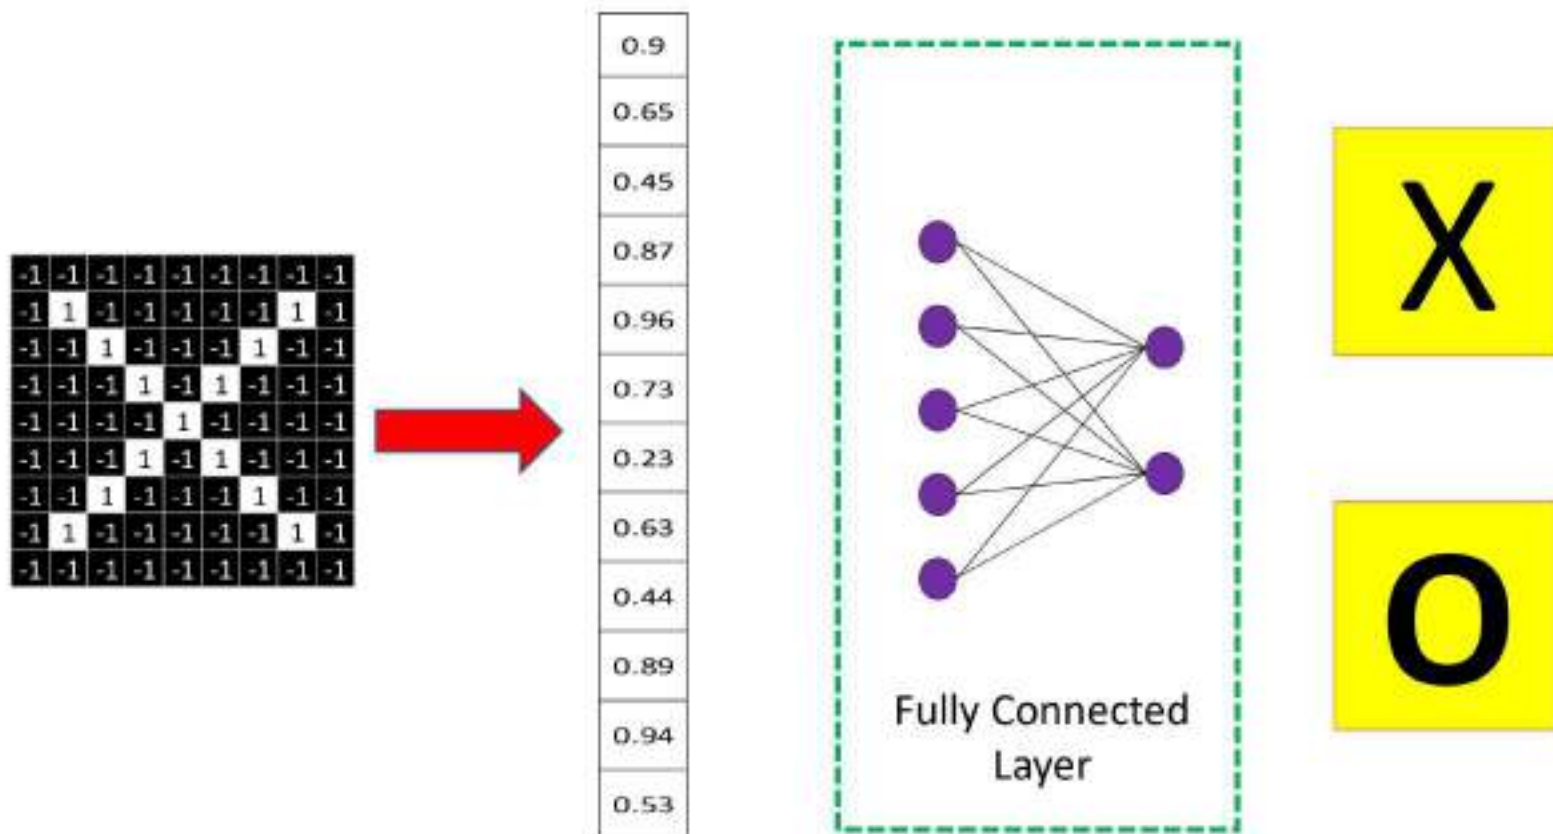
Fully Connected Layer (Testing Phase)



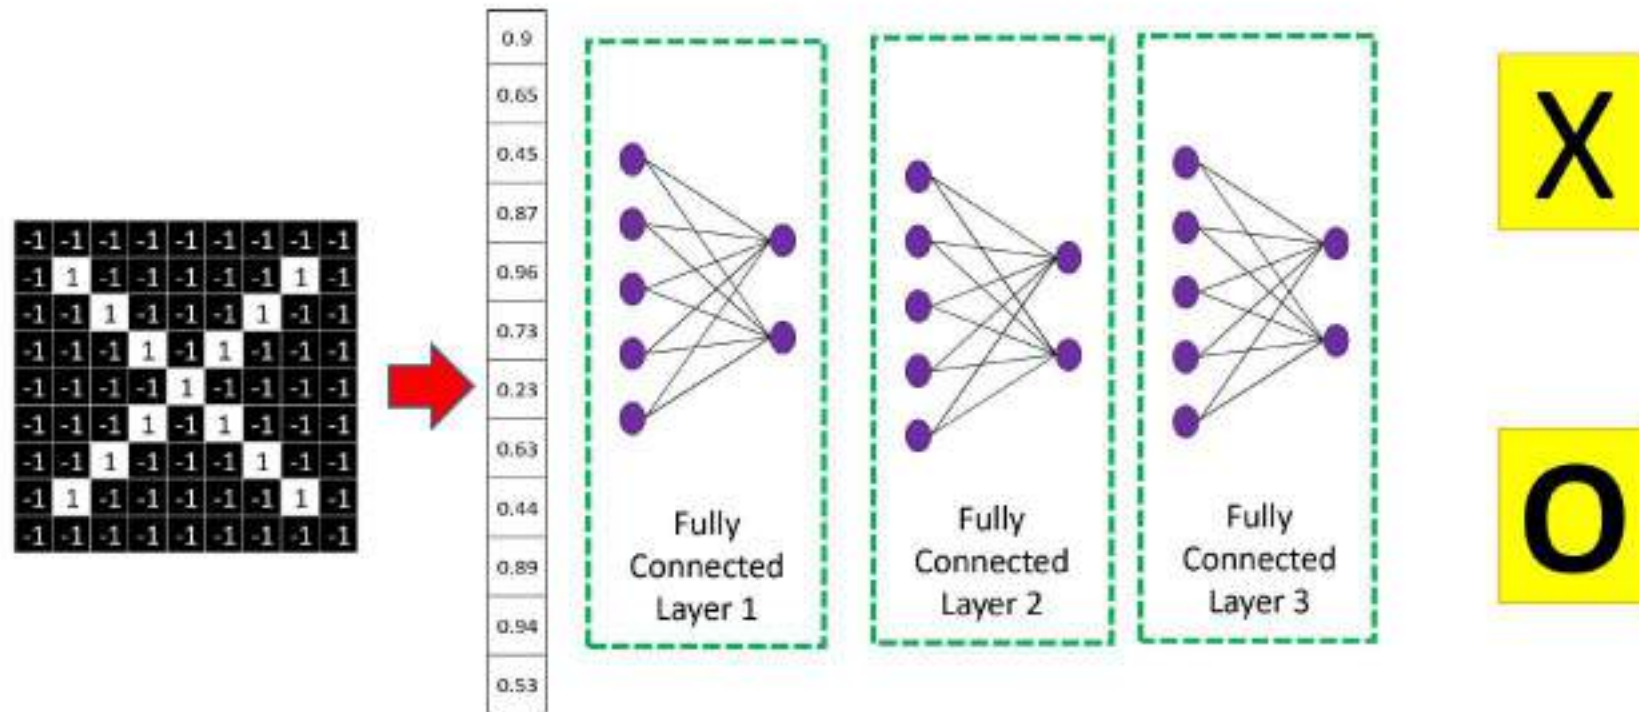
Fully Connected Layer (Testing Phase)



Fully Connected Layer

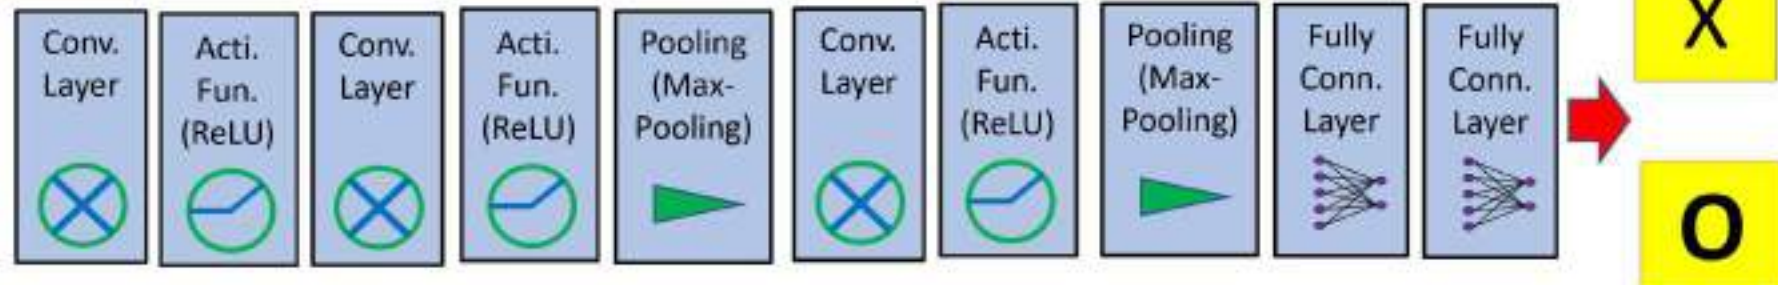


Multiple Stacking of Fully Connected Layers




Stacking of Multiple Layers

-1	-1	-1	-1	-1	-1	-1	-1	-1
-1	1	-1	-1	-1	-1	-1	1	-1
-1	-1	1	-1	-1	-1	1	-1	-1
-1	-1	-1	1	-1	1	-1	-1	-1
-1	1	-1	-1	1	-1	-1	-1	-1
-1	-1	-1	1	-1	1	-1	1	-1
-1	-1	1	-1	-1	-1	1	-1	-1
-1	1	-1	-1	-1	-1	-1	1	-1
-1	-1	-1	-1	-1	-1	-1	-1	-1



Convolutional Neural Network- Classification



- Convolutional layer and Pooling help to extract high level features of input
- Fully connected layer used extracted high level features for classification of input image in different classes
- Output also include the class probability of the image

Image Augmentation

To reduce the problem of overfitting, various augmentation variants can be applied to particular image of the training dataset.



Original Image



**90° left
rotation**



**90° right
rotation**



**180° Horizontal
rotation**



**Horizontal
Flip**



Vertical Flip



**Increased
Brightness**



**Addition of
noise**



Low contrast



**Background
Removal**

Applications of CNNs



**Face
Recognition**



**Analyzing
Documents**



**Medical
Imaging**



Applications of CNNs



Automatic Navigation

G.P. S.



Image
Recognition and
Decision making



Applications of CNNs

Image Recognition (*Prediction of objects in different classes*)

